



# **CSE488 (Section 1)** **[Spring 2023]**

## **Lab Assignment Submission Report**

### **All Online MapReduce LAB**

**Submitted by:**  
**Sazzad Hossen**  
**ID: 2019-1-60-063**

# LAB 1

1.

```
# 1. Find the Air craft, country flying mission and target country for those missions in which time on target was less than 1000.
```

```
query= """
```

```
SELECT TargetCountry, COUNT(TargetCountry)
FROM Bombing_Operations
WHERE TimeOnTarget < 1000
GROUP BY TargetCountry
```

```
"""
```

```
ass1=spark.sql(query)
ass1.show()
```

```
+-----+-----+
| TargetCountry|count(TargetCountry)|
+-----+-----+
| CAMBODIA|154639|
| WESTPAC WATERS|63|
| SOUTH VIETNAM|1578639|
| UNKNOWN|48|
| PHILLIPINES|1|
| THAILAND|53765|
| NORTH VIETNAM|432509|
| LAOS|703833|
+-----+-----+
```

2.

```
[39] #2. Find the target country and the number of times they have been attacked. Show the result in descending order.
```

```
query= """
```

```
SELECT MissionDate,TargetCountry
FROM Bombing_Operations B
ORDER BY B.TimeOnTarget DESC,
        B.MissionDate DESC
```

```
"""
```

```
ass2=spark.sql(query)
ass2.show()
```

```
+-----+-----+
|MissionDate|TargetCountry|
+-----+-----+
| 1965-11-25|NORTH VIETNAM|
| 1966-02-18|SOUTH VIETNAM|
| 1966-02-18|SOUTH VIETNAM|
| 1966-02-05|SOUTH VIETNAM|
| 1966-02-05|SOUTH VIETNAM|
| 1966-01-31|SOUTH VIETNAM|
| 1965-10-05|SOUTH VIETNAM|
| 1965-10-05|SOUTH VIETNAM|
| 1965-10-05|SOUTH VIETNAM|
| 1965-10-05|SOUTH VIETNAM|
| 1965-10-05|SOUTH VIETNAM|
+-----+-----+
```

3.

```
[55] # 3.Find the target country which had been attacked most.
      query= """
      SELECT TargetCountry,
             COUNT(*) TimeOnTarget
      FROM Bombing_Operations
      GROUP BY TargetCountry
      HAVING COUNT(TargetCountry) > 5
      """

      ass3=spark.sql(query)
      ass3.show()
```

TargetCountry	TimeOnTarget
CAMBODIA	218788
WESTPAC WATERS	74
SOUTH VIETNAM	2172403
UNKNOWN	59
THAILAND	60486
NORTH VIETNAM	694269
LAOS	1254695

4.

```
[57] # 4. Find the country flying mission which attacked most.
      query= """
      SELECT ContryFlyingMission,
             COUNT(*) TimeOnTarget
      FROM Bombing_Operations
      GROUP BY ContryFlyingMission
      HAVING COUNT(TimeOnTarget) > 1000
      """

      ass4=spark.sql(query)
      ass4.show()
```

ContryFlyingMission	TimeOnTarget
VIETNAM (SOUTH)	622013
KOREA (SOUTH)	24469
UNITED STATES OF ...	3708997
AUSTRALIA	12519
LAOS	32777

5.

[58] # 5. Find the take off location and how many times they had used.

```
query= """
SELECT TakeoffLocation,
       COUNT(TimeOnTarget) TimeOnTarget
FROM Bombing_Operations
GROUP BY TakeoffLocation

"""
```

```
ass5=spark.sql(query)
ass5.show()
```

TakeoffLocation	TimeOnTarget
LOUANG PHRAB	5164
TAN SON NHUT	311871
KHE SANH	27
NHA TRANG	56818
WESTPAC 59	481
KADENA AFB	5413
TONKIN GULF	194182
VINH BINH	10

6.

[87] # 6. Find the number of fighter jets from Aircraft\_Glossary.

```
Aircraft_Glossary.createOrReplaceTempView("Aircraft_Glossary")
query= """ SELECT COUNT(*) AirCraft FROM Aircraft_Glossary WHERE AirCraftType = 'Fighter Jet' """
```

```
ass6= spark.sql(query)
ass6.show()
```

AirCraft
10

7.

```
[98] # 7. Find the number of different types air craft for each air craft type.
```

```
query= """ SELECT DISTINCT Aircraft FROM Aircraft_Glossary """
```

```
ass7= spark.sql(query)
ass7.show()
```

```
+-----+
|Aircraft|
+-----+
|E-2|
|LC-130|
|F-5|
|OV-10|
|E-3|
|HC-47|
|HH-3|
|TA-4|
|EC-121|
|AC-123|
|C-76|
|OV-1|
|B-57|
|U-17|
|EB-66|
```

8.

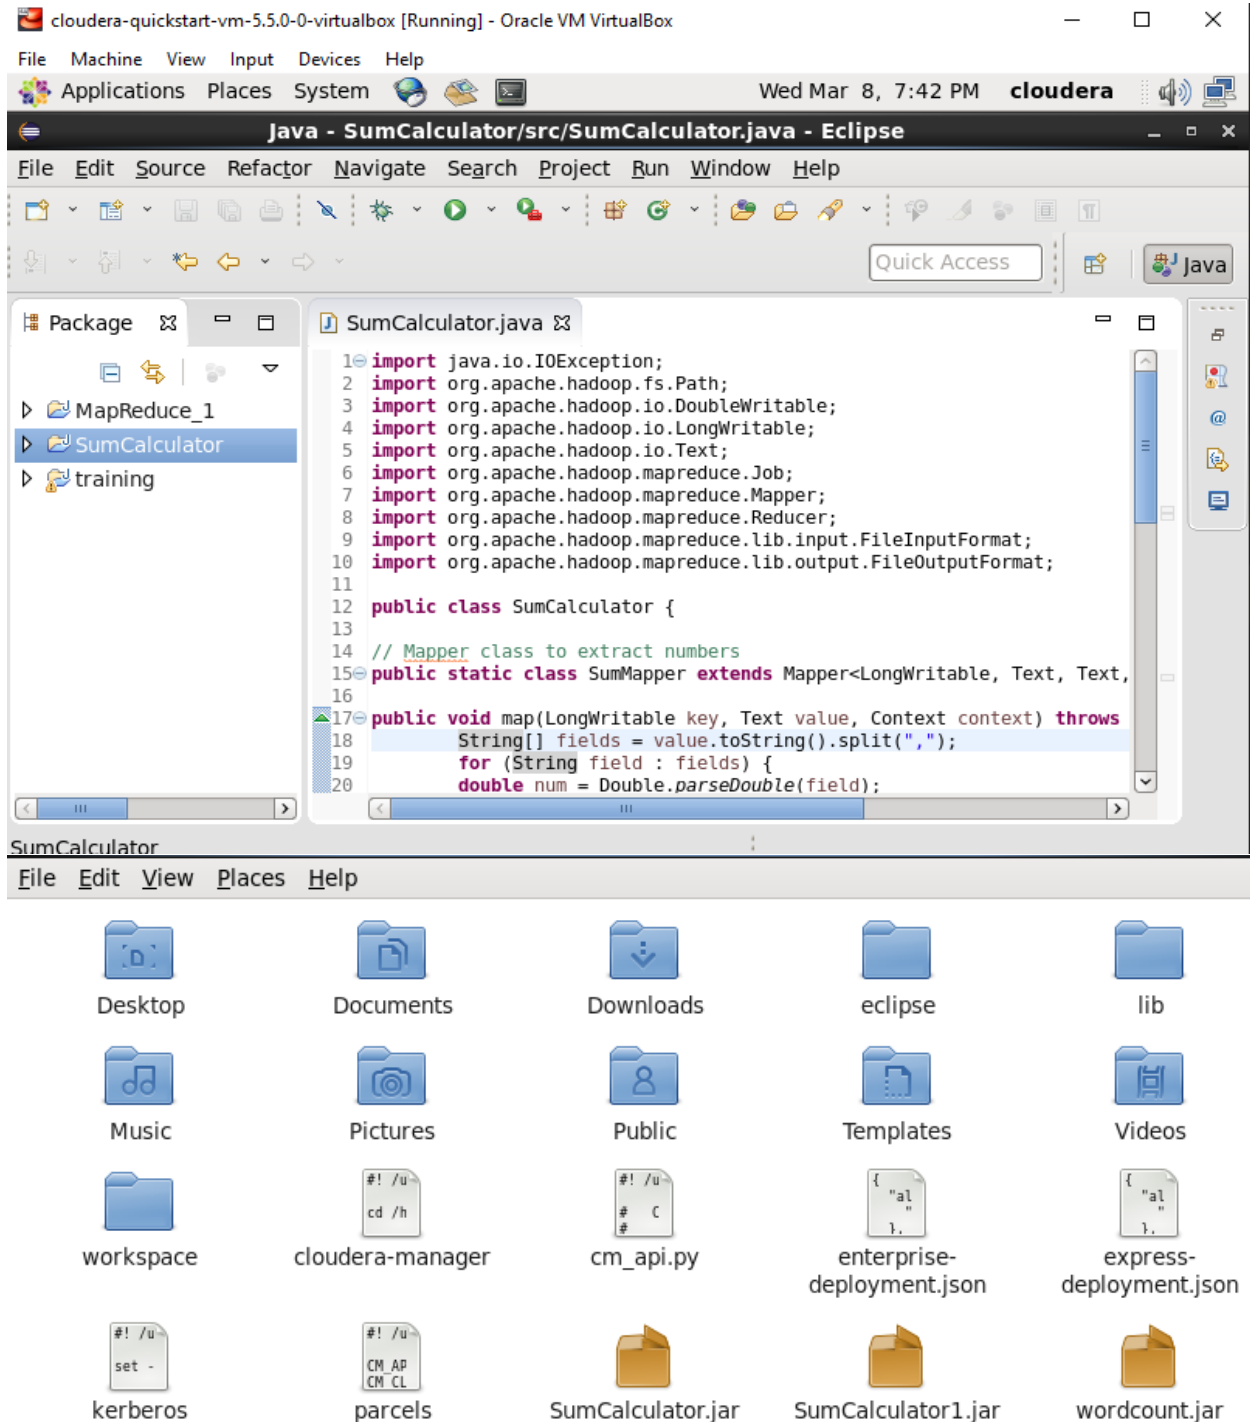
```
p5=Bombing_Operations.join(Aircraft_Glossary,Bombing_Operations.AirCraft == Aircraft_Glossary.Aircraft,"inner") \.show(truncate=False)
```

AirCraft	ContryFlyingMission	MissionDate	OperationSupported	PeriodOfDay	TakeoffLocation	TargetCountry	TimeOnTarget	WeaponType	WeaponsLoadedWeight
EC-47	UNITED STATES OF AMERICA	1971-06-05	null	D	TAN SON NHUT	CAMBODIA	1005.0	null	0
EC-47	UNITED STATES OF AMERICA	1972-12-26	null	D	NAKHON PHANOM	SOUTH VIETNAM	530.0	null	0
RF-4	UNITED STATES OF AMERICA	1973-07-28	null	D	UDORN AB	LAOS	730.0	null	0
A-1	UNITED STATES OF AMERICA	1970-02-02	null	N	NAKHON PHANOM	LAOS	1415.0	BLU27 FIRE BOMB (750)	17400
A-37	VIETNAM (SOUTH)	1970-10-08	null	D	DANANG	SOUTH VIETNAM	1240.0	null	0
F-4	UNITED STATES OF AMERICA	1970-11-25	null	D	UBON AB	LAOS	650.0	MK 82 GP BOMB (500) LD	31860
A-4	UNITED STATES OF AMERICA	1972-03-08	null	D	TONKIN GULF	LAOS	1005.0	null	0
F-4	UNITED STATES OF AMERICA	1971-12-27	null	null	UDORN AB	LAOS	0.0	null	0
A-7	UNITED STATES OF AMERICA	1972-05-24	null	null	TONKIN GULF	NORTH VIETNAM	0.0	null	0
EC-47	UNITED STATES OF AMERICA	1972-09-12	null	D	TAN SON NHUT	SOUTH VIETNAM	710.0	null	0
CH-53	UNITED STATES OF AMERICA	1974-06-13	null	N	NAKHON PHANOM	THAILAND	1800.0	null	0
CH-53	UNITED STATES OF AMERICA	1974-12-19	null	D	NAKHON PHANOM	THAILAND	800.0	null	0
O-1	VIETNAM (SOUTH)	1973-10-24	null	D	NHA TRANG	SOUTH VIETNAM	800.0	null	0
UH-1	VIETNAM (SOUTH)	1974-03-10	null	In	DMU CAT	SOUTH VIETNAM	1800.0	CARGO (TONS)	10

0s completed at 11:34 PM

# LAB 3

## Sum Calculator:



```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ hadoop jar SumCalculator1.jar SumCalculator /user/cloudera/wordcountfiles  
23/03/08 19:37:40 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032  
23/03/08 19:37:43 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed by Runner to remedy this.  
23/03/08 19:37:44 INFO input.FileInputFormat: Total input paths to process : 1  
23/03/08 19:37:44 INFO mapreduce.JobSubmitter: number of splits:1  
23/03/08 19:37:44 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1678328889585_0001  
23/03/08 19:37:45 INFO impl.YarnClientImpl: Submitted application application_1678328889585_0001  
23/03/08 19:37:45 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/pr  
23/03/08 19:37:45 INFO mapreduce.Job: Running job: job_1678328889585_0001  
23/03/08 19:38:15 INFO mapreduce.Job: Job job_1678328889585_0001 running in uber mode : false  
23/03/08 19:38:15 INFO mapreduce.Job: map 0% reduce 0%  
23/03/08 19:38:26 INFO mapreduce.Job: map 100% reduce 0%  
23/03/08 19:38:35 INFO mapreduce.Job: map 100% reduce 100%  
23/03/08 19:38:35 INFO mapreduce.Job: Job job_1678328889585_0001 completed successfully  
23/03/08 19:38:35 INFO mapreduce.Job: Counters: 49  
File System Counters  
FILE: Number of bytes read=90  
FILE: Number of bytes written=222989  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=146  
HDFS: Number of bytes written=15  
HDFS: Number of read operations=6  
HDFS: Number of large read operations=0  
Java - SumCalculator/src/SumCalculator.java - Eclipse hs=2
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Combine input records=0  
Combine output records=0  
Reduce input groups=1  
Reduce shuffle bytes=90  
Reduce input records=6  
Reduce output records=1  
Spilled Records=12  
Shuffled Maps =1  
Failed Shuffles=0  
Merged Map outputs=1  
GC time elapsed (ms)=214  
CPU time spent (ms)=1360  
Physical memory (bytes) snapshot=330555392  
Virtual memory (bytes) snapshot=3008303104  
Total committed heap usage (bytes)=226365440  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Input Format Counters  
Bytes Read=11  
File Output Format Counters  
Bytes Written=15  
Hue - File Browser - Mozilla Firefox  
Hue - ... Java - ... Comp... / [home] cloudera cloude...
```

Hue - File Browser - Mozilla Firefox

about:s...estore Hue - Fi... x East West ... CSE488 [S... CSE488\_L... Google Ac... +

quickstart.cloudera:8888/filebrowser/view=/user/ Search

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager

Query Editors Data Browsers Workflows Search Security

File Browser

Search for file name

Actions Move to trash

	Name	Size	User	Group	Permissions	Date
			cloudera	cloudera	drwxr-xr-x	March 08, 2023 07:38 PM
	.		cloudera	cloudera	drwxr-xr-x	March 08, 2023 07:38 PM
	_SUCCESS	0 bytes	cloudera	cloudera	-rw-r--r--	March 08, 2023 07:38 PM
	part-r-00000	15 bytes	cloudera	cloudera	-rw-r--r--	March 08, 2023 07:38 PM

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager

Query Editors Data Browsers Workflows Search Security

File Browser

ACTIONS

View as binary

Edit file

Download

Home

Page 1 of 1

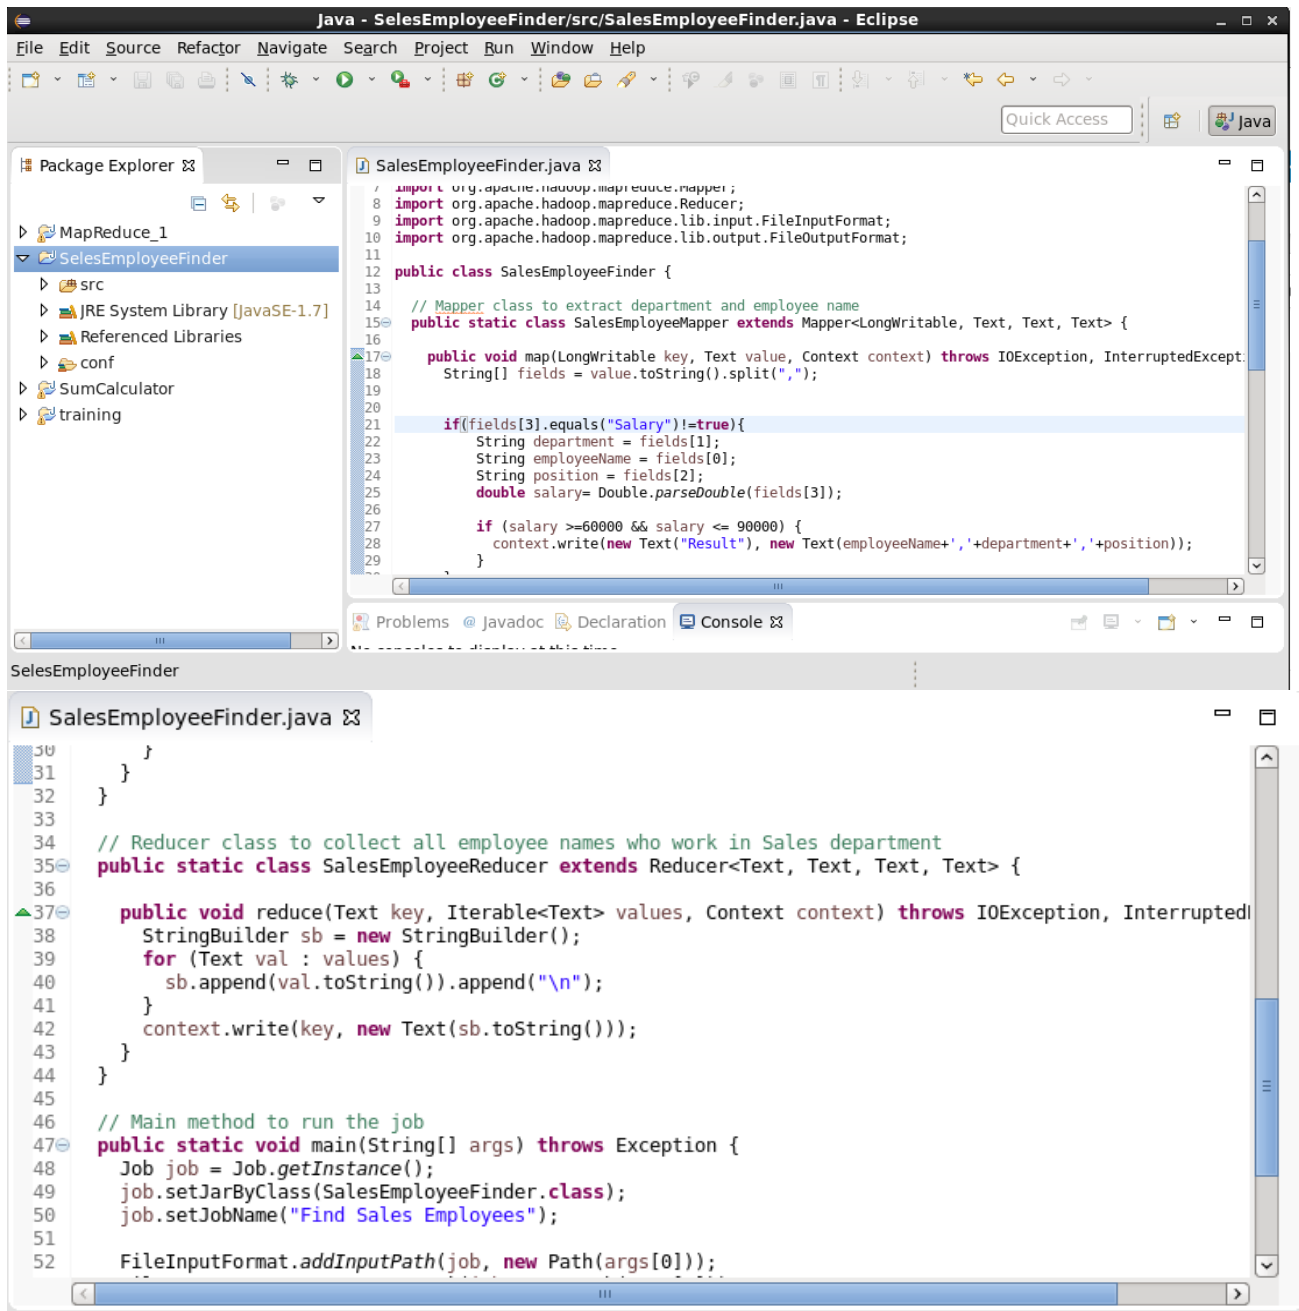
/ user / cloudera / output / part-r-00000

Total sum 23.0



# LAB 4

## INPUT:



```
Java - SelesEmployeeFinder/src/SalesEmployeeFinder.java - Eclipse
File Edit Source Refactor Navigate Search Project Run Window Help

Package Explorer
MapReduce_1
SelesEmployeeFinder
src
JRE System Library [JavaSE-1.7]
Referenced Libraries
conf
SumCalculator
training

SalesEmployeeFinder.java
7 import org.apache.hadoop.mapreduce.Mapper;
8 import org.apache.hadoop.mapreduce.Reducer;
9 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
10 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
11
12 public class SalesEmployeeFinder {
13
14     // Mapper class to extract department and employee name
15     public static class SalesEmployeeMapper extends Mapper<LongWritable, Text, Text, Text> {
16
17         public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
18             String[] fields = value.toString().split(",");
19
20             if(fields[3].equals("Salary")!=true){
21                 String department = fields[1];
22                 String employeeName = fields[0];
23                 String position = fields[2];
24                 double salary = Double.parseDouble(fields[3]);
25
26                 if (salary >= 60000 && salary <= 90000) {
27                     context.write(new Text("Result"), new Text(employeeName+' '+department+' '+position));
28                 }
29             }
30         }
31     }
32 }
33
34 // Reducer class to collect all employee names who work in Sales department
35 public static class SalesEmployeeReducer extends Reducer<Text, Text, Text, Text> {
36
37     public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {
38         StringBuilder sb = new StringBuilder();
39         for (Text val : values) {
40             sb.append(val.toString()).append("\n");
41         }
42         context.write(key, new Text(sb.toString()));
43     }
44 }
45
46 // Main method to run the job
47 public static void main(String[] args) throws Exception {
48     Job job = Job.getInstance();
49     job.setJarByClass(SalesEmployeeFinder.class);
50     job.setJobName("Find Sales Employees");
51
52     FileInputFormat.addInputPath(job, new Path(args[0]));
53 }
```



Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started

HUE Home Query Editors Data Browsers Workflows Search Security File Browser

### File Browser

**ACTIONS**

- View as binary
- Edit file
- Download
- View file location
- Refresh

**INFO**

Last modified: March 15, 2023 7:51 p.m.

Home / user / cloudera / seles / sample\_data.csv

```
Name,Department,Position,Salary,Start_Date
John Smith,Engineering,Software Developer,75000,1/15/2019
Emily Johnson,Marketing,Marketing Coordinator,55000,3/12/2018
Michael Lee,Finance,Financial Analyst,65000,7/1/2020
Sarah Chen,Human Resources,HR Manager,85000,9/5/2017
David Kim,Operations,Operations Manager,95000,4/2/2016
Jasmine Wong,Sales,Sales Representative,60000,2/21/2021
Matthew Johnson,Engineering,Engineering Manager,100000,11/10/2015
Ava Martinez,Marketing,Marketing Manager,90000,6/23/2014
Jessica Park,Finance,Financial Controller,110000,2/15/2013
Richard Lee,Operations,Operations Coordinator,55000,1/1/2020
```

### Output:

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started

HUE Home Query Editors Data Browsers Workflows Search Security File Browser

### File Browser

**ACTIONS**

- View as binary
- Edit file
- Download
- View file location
- Refresh

**INFO**

Last modified

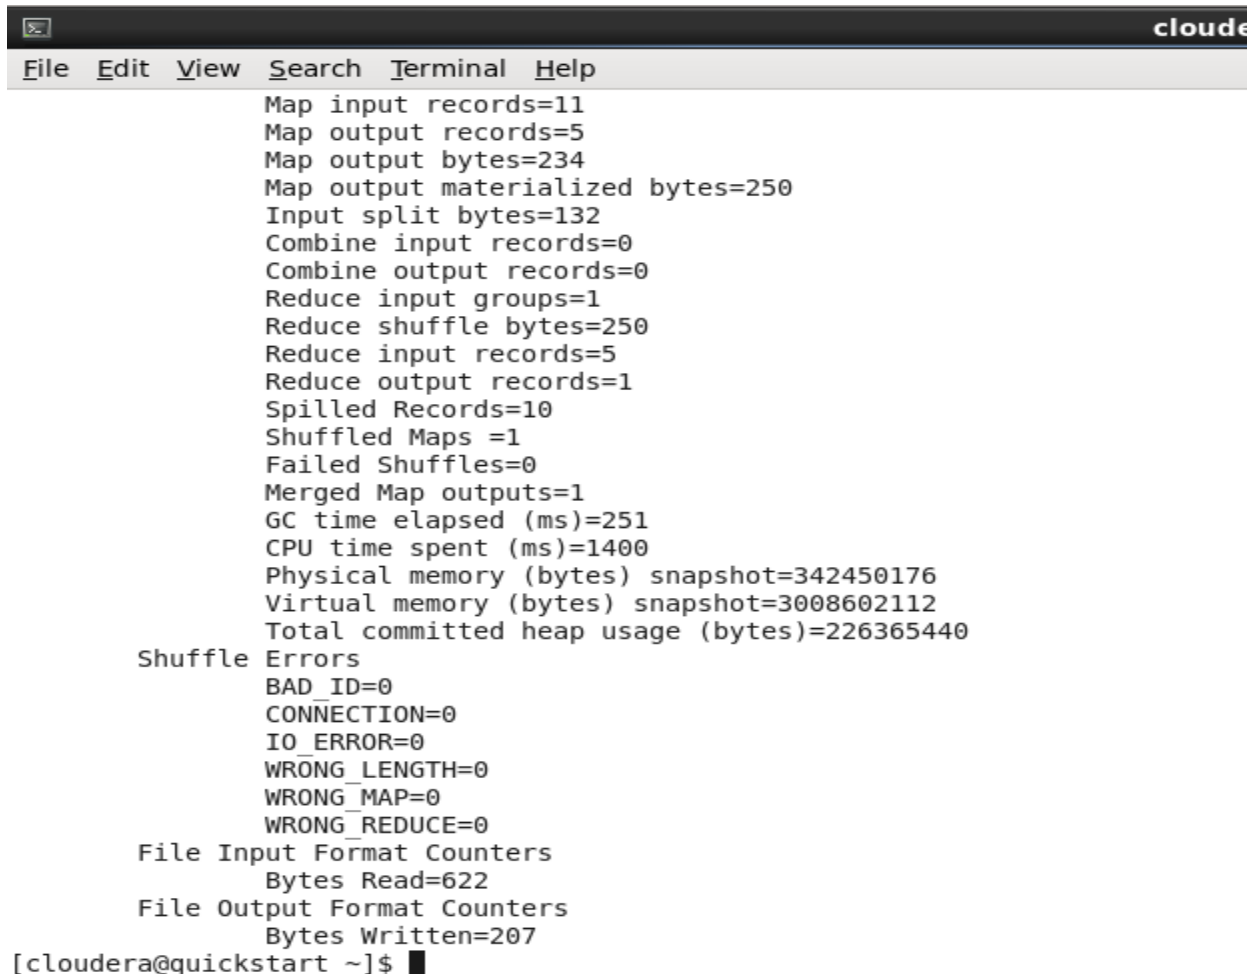
Home / user / cloudera / selesoutput2 / part-r-00000

```
Result Ava Martinez,Marketing,Marketing Manager
Jasmine Wong,Sales,Sales Representative
Sarah Chen,Human Resources,HR Manager
Michael Lee,Finance,Financial Analyst
John Smith,Engineering,Software Developer
```

```

[cloudera@quickstart ~]$ hadoop jar SelesEmployeeFinder1.jar SalesEmployeeFinder
/user/cloudera/seles /user/cloudera/selesoutput2
23/03/15 20:32:05 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0
:8032
23/03/15 20:32:06 WARN mapreduce.JobResourceUploader: Hadoop command-line option
 parsing not performed. Implement the Tool interface and execute your applicatio
n with ToolRunner to remedy this.
23/03/15 20:32:07 INFO input.FileInputFormat: Total input paths to process : 1
23/03/15 20:32:07 INFO mapreduce.JobSubmitter: number of splits:1
23/03/15 20:32:07 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_16
78934192891_0003
23/03/15 20:32:07 INFO impl.YarnClientImpl: Submitted application application_16
78934192891_0003
23/03/15 20:32:07 INFO mapreduce.Job: The url to track the job: http://quickstar
t.cloudera:8088/proxy/application_1678934192891_0003/
23/03/15 20:32:07 INFO mapreduce.Job: Running job: job_1678934192891_0003
23/03/15 20:32:21 INFO mapreduce.Job: Job job_1678934192891_0003 running in ub
er mode : false
23/03/15 20:32:21 INFO mapreduce.Job: map 0% reduce 0%
23/03/15 20:32:40 INFO mapreduce.Job: map 100% reduce 0%
23/03/15 20:32:53 INFO mapreduce.Job: map 100% reduce 100%
23/03/15 20:32:53 INFO mapreduce.Job: Job job_1678934192891_0003 completed suc
cessfully
23/03/15 20:32:53 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=250
        FILE: Number of bytes written=223361
        FILE: Number of read operations=0

```



```

File Edit View Search Terminal Help
Map input records=11
Map output records=5
Map output bytes=234
Map output materialized bytes=250
Input split bytes=132
Combine input records=0
Combine output records=0
Reduce input groups=1
Reduce shuffle bytes=250
Reduce input records=5
Reduce output records=1
Spilled Records=10
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=251
CPU time spent (ms)=1400
Physical memory (bytes) snapshot=342450176
Virtual memory (bytes) snapshot=3008602112
Total committed heap usage (bytes)=226365440
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=622
File Output Format Counters
    Bytes Written=207
[cloudera@quickstart ~]$ █

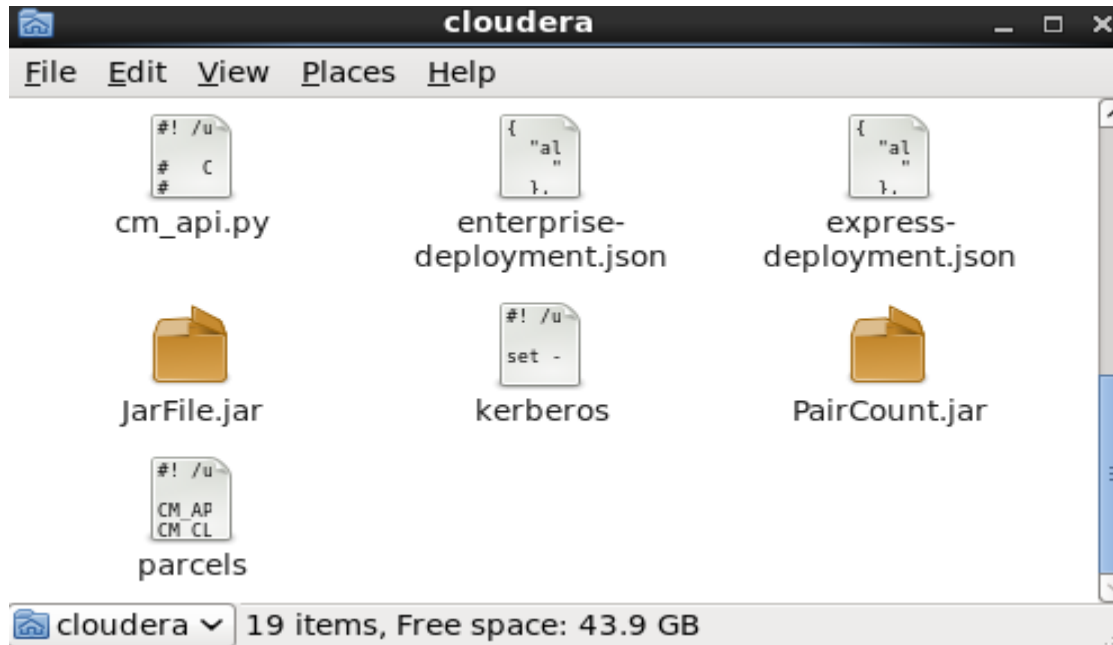
```

# LAB 5


## JAVA FILE:

```
PairCount.java
22     private Text p = new Text();
23
24     public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
25         String line = value.toString();
26         StringTokenizer tokenizer = new StringTokenizer(line);
27         String[] items = tokenizer.nextToken().split(",");
28
29         Arrays.sort(items);
30
31         for (int i = 0; i < items.length - 1; i++) {
32             for (int j = i + 1; j < items.length; j++) {
33
34                 if(items[i].charAt(0)<items[j].charAt(0)){
35
36                     p.set(items[i] + "," + items[j]);
37                     context.write(p, o);
38                 }
39                 else{
40                     p.set(items[j] + "," + items[i]);
41                     context.write(p, o);
42                 }
43             }
44         }
45     }
46 }
47
48
49     public static class PairCountReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
50
51         private IntWritable count = new IntWritable();
52
53         public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {
54             int sum = 0;
55
56             Iterator<IntWritable> iterator = values.iterator();
57
58             while (iterator.hasNext()) {
59                 sum += iterator.next().get();
60             }
61             count.set(sum);
62             context.write(key, count);
63         }
64     }
65
66     public static void main(String[] args) throws Exception {
67         Configuration conf = new Configuration();
68         Job job = Job.getInstance(conf, "pair count");
69     }
```

### JAR file Create “Pair Count”:



### Dataset input in Hue Cloudera:



## Run in Terminal:

```
cloudera@quickstart:~$  
File Edit View Search Terminal Help  
Bytes Written=78  
[cloudera@quickstart ~]$ hadoop jar PairCount.jar PairCount /user/cloudera/paircount /user/cloudera/paircountoutput6  
23/03/22 20:22:03 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032  
23/03/22 20:22:04 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.  
23/03/22 20:22:04 INFO input.FileInputFormat: Total input paths to process : 1  
23/03/22 20:22:04 INFO mapreduce.JobSubmitter: number of splits:1  
23/03/22 20:22:04 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1677120724034_0009  
23/03/22 20:22:05 INFO impl.YarnClientImpl: Submitted application application_1677120724034_0009  
23/03/22 20:22:05 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1677120724034_0009/  
23/03/22 20:22:05 INFO mapreduce.Job: Running job: job_1677120724034_0009  
23/03/22 20:22:14 INFO mapreduce.Job: Job job_1677120724034_0009 running in uber mode : false  
23/03/22 20:22:14 INFO mapreduce.Job:  map 0% reduce 0%  
23/03/22 20:22:23 INFO mapreduce.Job:  map 100% reduce 0%  
23/03/22 20:22:32 INFO mapreduce.Job:  map 100% reduce 100%  
23/03/22 20:22:33 INFO mapreduce.Job: Job job_1677120724034_0009 completed successfully
```

```
File Edit View Search Terminal Help
Reduce shuffle bytes=136
Reduce input records=13
Reduce output records=13
Spilled Records=26
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=181
CPU time spent (ms)=1350
Physical memory (bytes) snapshot=326934528
Virtual memory (bytes) snapshot=3008466944
Total committed heap usage (bytes)=226365440
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=35
File Output Format Counters
  Bytes Written=78
[cloudera@quickstart ~]$
```

Output:

Query Editors ▾ Data Browsers ▾ Workflows ▾ Search Security ▾

File Browser

ACTIONS

View as binary

Edit file

Download

View file location

Refresh

INFO

Last modified  
March 22, 2023 8:22 p.m.

Home / [user](#) / [cloudera](#) / [paircountoutput6](#) / [part-r-00000](#)

A, B	4
A, C	2
A, D	3
A, E	3
A, K	1
B, C	2
B, D	3
B, E	2
B, K	1
C, D	1
C, E	2
D, E	1